


Nonlinear Regression



Prof. Fabrício Olivetti de França

Federal University of ABC

05 February, 2024



Nonlinear Regression

We now turn our attention to a more generic regression model

$$y^{(i)} = f(x^{(i)}; \theta) + \epsilon^{(i)}$$

As before, $x^{(i)}$ is the vector with the regressor variables, ϵ is a normally distributed disturbance, and f is the expectation function of a parameterized family of functions specified by the parameters θ .

We use θ for this function to differentiate from the linear models where the parameters were labeled β .

A **nonlinear regression** model is characterized by the function $f(x; \theta)$ where at least one of the derivatives of the expectation function wrt the parameters depends on at least one of the parameters.

Example

For example, the models:

$$f(x; \theta) = 60 + 70e^{-\theta x}$$

$$f(x; \theta) = \frac{\theta_1 x}{\theta_2 + x}$$

are both nonlinear.

Example

All of the partial derivatives depend on at least one parameter:

$$\begin{aligned}\frac{\partial f(x; \theta)}{\partial \theta} &= -70xe^{-\theta x} \\ \frac{\partial f(x; \theta)}{\partial \theta_1} &= \frac{x}{\theta_2 + x} \\ \frac{\partial f(x; \theta)}{\partial \theta_2} &= \frac{-\theta_1 x}{(\theta_2 + x)^2}\end{aligned}$$

Fitting nonlinear models requires nonlinear optimization methods that do not guarantee convergence to the global optima.

We can resort to linear approximations to fit the parameters.

Transformably Linear Models

Notice that our second example can be changed into a linear model:

$$\begin{aligned}f(x; \theta) &= \frac{\theta_1 x}{\theta_2 + x} \\ \frac{1}{f(x; \theta)} &= \frac{\theta_2 + x}{\theta_1 x} \\ \frac{1}{f(x; \theta)} &= \frac{\theta_2}{\theta_1 x} + \frac{x}{\theta_1 x} \\ \frac{1}{f(x; \theta)} &= \frac{1}{\theta_1} + \frac{\theta_2}{\theta_1 x} \\ g(u; \beta) &= \beta_1 + \beta_2 u \\ u &= \frac{1}{x}\end{aligned}$$

This is called **transformably linear models**.

Notice, though, that when we transform f to a linear model, we are also transforming the disturbance term.

This can change the distribution of the disturbance, making the assumptions invalid.

This linearization can be helpful to find initial guesses for the parameters values, but the fitting should be performed with a nonlinear optimization method.

Conditionally Linear Parameters

Some parameters of a nonlinear model may be **conditionally linear**. For example, the derivative wrt θ_1 of:

$$f(x; \theta) = \frac{\theta_1 x}{\theta_2 + x}$$
$$\frac{\partial f(x; \theta)}{\partial \theta_1} = \frac{x}{\theta_2 + x}$$

depends only on θ_2 . We can exploit this by fitting θ_1 as a linear regression conditional to a value of θ_2 .

Adjusting the parameters

Least Squares Estimate

As before, we want to find $\hat{\theta}$ such that:

$$\hat{\theta} = \min_{\theta} \|y - \hat{y}\|^2$$

$$\hat{y} = f(x; \theta)$$

The **Gauss-Newton** method starts with the idea of using a linear approximation of the expectation function. This can be done by Taylor expansion at an initial guess $\theta^{(0)}$:

$$f(x; \theta) = f(x; \theta^{(0)}) + \sum_{i=1}^p \frac{\partial f(x; \theta)}{\partial \theta_i} \Big|_{\theta=\theta^{(0)}} (\theta_i - \theta_i^{(0)})$$

The **Jacobian matrix** (J) is the matrix where:

$$J = \begin{bmatrix} \left. \frac{\partial f(x^{(1)}; \theta)}{\partial \theta_1} \right|_{\theta=\theta^{(0)}} & \cdots & \left. \frac{\partial f(x^{(1)}; \theta)}{\partial \theta_P} \right|_{\theta=\theta^{(0)}} \\ \vdots & \ddots & \vdots \\ \left. \frac{\partial f(x^{(N)}; \theta)}{\partial \theta_1} \right|_{\theta=\theta^{(0)}} & \cdots & \left. \frac{\partial f(x^{(N)}; \theta)}{\partial \theta_P} \right|_{\theta=\theta^{(0)}} \end{bmatrix}$$

So we have:

$$f(x; \theta) \approx f(x; \theta^{(0)}) + J \cdot (\theta - \theta^{(0)})^T$$

with J as a $N \times p$ matrix and θ a $1 \times p$ vector.

The approximated residuals are:

$$z(\theta) \approx y - (f(x; \theta^{(0)}) + J \cdot (\theta - \theta^{(0)}))$$

$$z(\theta) \approx (y - (f(x; \theta^{(0)}))) - J \cdot (\theta - \theta^{(0)})$$

$$z(\theta) \approx z^{(0)} - J\delta$$

Now, we want to find $\delta^{(0)}$ that minimizes $\|z^{(0)} - J^{(0)}\delta\|^2$.

This relates to our solution of least squares, where we wanted to minimize $\|y - X\beta\|^2$.

$$\begin{aligned}J^{(0)} &= QR = Q_1 R_1 \\ \delta^{(0)} &= R_1^{-1} Q_1^T z^{(0)}\end{aligned}$$

the point $f(x; \theta^{(0)} + \delta^{(0)})$ should be closer to y than $f(x; \theta^{(0)})$.

This process is repeated until we cannot minimize the sum of squared residues.

```
1 df = pd.read_csv("grade.csv")
2 xcols = ['ETA_mean']
3
4 x, y = df[xcols].values, df.grade.values[:,np.newaxis]
5
6 def f(x, theta):
7     return theta[0]*x/(theta[1] + x)
8 def J(x, theta):
9     return np.hstack([x/(theta[1]+x),
10                       -theta[0]*x/(theta[1]+x)**2])
11 def sumofsquares(x, y, t, f):
12     return np.sum(np.square(y - f(x, t)))
```

```
1 def gauss_step(x, y, t0, f, J):
2     p = t0.shape[0]
3     z = y - f(x, t0)
4     q, r = sc.linalg.qr(J(x, t0))
5     q1 = q[:, :p]
6     r1 = r[:p, :p]
7     delta = sc.linalg.solve_triangular(r1, q1.T@z)
8     return t0 + delta
```

Gauss-Newton Method

```
1 def gauss(x, y, t, f, J, debug=False):
2     sos_old = np.inf
3     sos = sumofsquares(x, y, t, f)
4     if debug:
5         print(sos)
6     while sos < sos_old:
7         t = gauss_step(x,y,t,f,J)
8         sos_old = sos
9         sos = sumofsquares(x, y, t, f)
10        if debug:
11            print(sos)
12    return t
```

```
1 t = np.array([205, 0.08])[:,np.newaxis]  
2 gauss(x, y, t, f, J, True)
```

245431183.92754266

132435.0246901058

12299140970.505436

array([[3.1590888],
 [-0.77757744]])

The last step increased the sum of squares even though we went towards the descent direction.

In these situations we need to decrease the step size when updating the current value of θ .

A simple heuristic is to start with a step size of 1 and half the current step size until the value becomes smaller than a threshold.

A better approach is to perform a line search to find the optimal value for the step size.

Step factor

```
1 def gauss_step(x, y, t0, f, J, step):
2     p = t0.shape[0]
3     z = y - f(x, t0)
4     q, r = sc.linalg.qr(J(x, t0))
5     q1 = q[:, :p]
6     r1 = r[:p, :p]
7     delta = sc.linalg.solve_triangular(r1, q1.T@z)
8     return t0 + step*delta
```

Step factor

```
1 def gauss(x, y, t, f, J, debug=False):
2     sos_old = np.inf
3     sos = sumofsquares(x, y, t, f)
4     step = 1
5     while step > 1e-12
6         and np.square(sos - sos_old) > 1e-6:
7             t_i = gauss_step(x,y,t,f,J,step)
8             sos_old = sos
9             sos = sumofsquares(x, y, t_i, f)
10            if sos > sos_old:
11                step = step / 2
12            else:
13                t = t_i
14    return t
```

```
1 t = np.array([205, 0.08])[:,np.newaxis]  
2 gauss(x, y, t, f, J, True)
```

Step factor

```
245431183.92754266
132435.0246901058
12299140970.505436
745214728.2425426
186876895.48195535
47131373.71276204
...
335857.52272309375
335857.5118304746
335857.5091074513
335857.50842667697
array([[ 0.00344322],
       [-0.35566967]])
```

Let's try a different starting value:

```
1 t = np.array([1.2, -0.1])[:,np.newaxis]
2 gauss(x, y, t, f, J, True)
```

```
208543.89714250065
28984.677830013265
28918.85467418468
28908.738157793916
28907.728409941887
28907.546226136466
28907.519349759037
28907.515040983937
28907.51437219738
array([[ 5.10148992],
       [-0.09892466]])
```

Choosing different starting points can lead to different optima, better or worse than the previous attempt.

Confidence Intervals

Approximating the Confidence Interval

With the linear models we could find the joint $(1-\alpha)\%$ CI of the parameters with:

$$(\beta - \hat{\beta})^T x^T x (\beta - \hat{\beta}) \leq P s^2 F(\alpha, P, N - P)$$

Approximating the Confidence Interval

We can do the same with the linear approximation of the nonlinear model:

$$(\theta - \hat{\theta})^T J^T J(\theta - \hat{\theta}) \leq P s^2 F(\alpha, P, N - P)$$

And generate the ellipse with:

$$\left\{ \theta = \hat{\theta} + \sqrt{Ps^2 F(\alpha, P, N - P)} R_1^T d \mid \|d\| = 1 \right\}$$

Similarly, the marginal CI can be calculated as:

$$s^2 = \frac{(y - f(x; \hat{\theta}))^2}{N - P}$$
$$se = \sqrt{s^2 \text{diag}(J^T J)}$$
$$\hat{\theta} \pm se \cdot t(\alpha/2, N - P)$$

Approximating the Confidence Interval

```
1 j = J(x, theta)
2 s2 = np.sum((y - f(x,theta))**2)/(n-p)
3 se = np.sqrt(s2 * np.diag(np.linalg.inv(j.T @ j)))
4
5 q, r = sc.linalg.qr(j)
6 q1 = q[:, :p]
7 r1 = r[:p, :p]
```

Approximating the Confidence Interval

```
1 alpha = 0.05
2 marginal_t = []
3 for i, coef in enumerate(theta.flatten()):
4     v = np.abs(se[i]*stats.t.ppf(1 - alpha/2, n-p))
5     marginal_t.append((coef - v, coef + v))
6     print(f"{coef - v:0.4f} <= theta_{i} <= {coef + v:0.4f}")
```


Approximating the Confidence Interval

```
1 omegas = np.pi*np.arange(0, 2.01, 0.01)
2 const = np.sqrt(p*s2*stats.f.ppf(1-alpha, p, n-p))
3         *np.linalg.inv(r1)
4 thetapts = np.array([theta + const
5                       @ np.array([np.cos(w), np.sin(w)])
6                       for w in omegas])
```


Approximating the Confidence Interval

```
1 _,ax = plt.subplots()
2 ax.plot(thetapts[:,0], thetapts[:,1], color='black')
3 ax.plot(theta[0], theta[1], '+', color='black')
4 ax.axvline(x=marginal_t[0][0], linestyle='--', color='black')
5 ax.axvline(x=marginal_t[0][1], linestyle='--', color='black')
6 ax.axhline(y=marginal_t[1][0], linestyle='--', color='black')
7 ax.axhline(y=marginal_t[1][1], linestyle='--', color='black')
8
9 ax.set_xlabel(r"$\theta_1$")
10 ax.set_ylabel(r"$\theta_2$")
```

Approximating the Confidence Interval



Approximating the Confidence Interval

We can calculate the CI of the predictions with:

$$\hat{y}^{(i)} \pm s \|J_i R_1^{-1}\| t(\alpha/2, N - P)$$
$$\hat{y} \pm s \|J R_1^{-1}\| \sqrt{f(\alpha, P, N - P)}$$

Approximating the Confidence Interval

```
1  for i, y_i in enumerate(f(x, theta).flatten()):
2      jr = (j[0,:] @ np.linalg.inv(r1))
3      v = np.sqrt(s2 * (jr.T@jr))*np.abs(stats.t.ppf(1 - alpha/2, n-p))
4      print(f"{y_i - v:0.2} <= y_{i} <= {y_i + v:0.2}")
```


Approximating the Confidence Interval

```
1 fxi = np.arange(0.1, 0.8, 0.05)[: , np.newaxis]
2 j = J(xi, theta)
3
4 yy = np.linalg.norm(j @ np.linalg.inv(r1), axis=1)
5      * np.sqrt(s2*p * stats.f.ppf(1-alpha, p, n-p))
6 y_h = f(xi, theta)
```

Approximating the Confidence Interval

```
1 _,ax = plt.subplots()
2 ax.plot(xi[:,0], y_h, '.', color='black', markersize=15)
3 ax.fill_between(xi[:,0], y_h - yy, y_h + yy, edgecolor='gray', facecolor='gray',
4                 linewidth=4, linestyle='dashdot', antialiased=True)
5 ax.set_xlabel('x_1')
6 ax.set_ylabel(r'$\hat{y}$')
```

Approximating the Confidence Interval



When we have more than 2 parameters, we can plot pairwise approximations of the joint intervals by using the i, j -th rows and columns of R_1 .


Approximating the Confidence Interval

```
1 def f(x, theta):  
2     return np.exp(-theta[0]*x) + theta[1]*x/(theta[2] + x)  
3 def J(x, theta):  
4     return np.hstack([-x*np.exp(-theta[0]*x),  
5                       x/(theta[1]+x), -theta[0]*x/(theta[1]+x)**2])  
6  
7 theta2 = np.array([0.1, 0.1, 0.1])[:,np.newaxis]  
8 theta2 = gauss(x, y, theta2, f, J)
```

Approximating the Confidence Interval

```
1  omegas = np.pi*np.arange(0, 2.01, 0.01)
2
3  _,axs = plt.subplots(2,2, figsize=(12,12))
4  for i in range(2):
5      for j in range(i+1, 3):
6          ixs = [i,j]
7          const = np.sqrt(p*s2*stats.f.ppf(1-alpha, p, n-p))
8                  *np.linalg.inv(r1[[i,j], :][:,[i,j]])
9          thetapts = np.array([theta2[ixs].flatten()
10                               + const
11                               @ np.array([np.cos(w), np.sin(w)]) for w in omegas])
```

Approximating the Confidence Interval



Closing

- **nonlinear regression:** is a model where at least one of the derivatives of the expectation function wrt the parameters depends on at least one of the parameters
- **transformably linear models:** are nonlinear models that can be transformed to linear.
- **conditionally linear:** are models that some of the parameters are linear when we fix the other parameter values.

- Chapter 2 and 3 of Bates, Douglas. “Nonlinear regression analysis and its applications.” Wiley Series in Probability and Statistics (1988).

- Symbolic Regression



Acknowledgments